

Social Network Analysis of the Enron Corpus

Richard Careaga

April 25, 2019

Introduction

Goal

The goal of this paper is to illustrate techniques of social network analysis in combination with natural language processing to identify discrete email subsets in the Enron Corpus.¹

The Enron Corpus is a collection of 500,000 emails obtained by the Federal Energy Regulatory Commission in plaintext form for a regulatory investigation, made public pursuant to a Freedom of Information Act request. In major litigation, it is not unusual for comparable volumes of email to be collected and reviewed. The conventional method of review is a keyword search.² Inevitably, the large majority of emails are barren of useful information.

Method

The process can be improved by preprocessing emails to be reviewed to construct an internal social network through methods of graph analysis. For this paper a latent cluster random effects model was applied.³

To provide an informal test of the efficacy of the latent graph classification, the vocabulary of each group was compared. Each shared distinct words in common, but each had unshared distinct terms with other groups. One group had a vocabulary with approximately 23.01% distinct words that did not appear in either other cluster.

As a method of reviewing emails, the machine learning approach of this approach has two principal benefits. The minimum information needed, a unique identifier for sender and receiver is either already available or extracted early in the process and so represents no additional effort. Judicious subsetting of users, based on graph metrics of centrality, to reduce the graph size, reduces the most computationally intensive portion of the work, latent model fitting. The second benefit is the ability to prioritize review of emails by graph cluster and with knowledge of the relative positions of the participants in the social network of the organization.

Background

In times of political turmoil, events can move from impossible to inevitable without even passing through improbable. [Anatole Kalesky](#)

[Enron Corp.](#) and its affiliates were engaged in energy-related businesses, as described in its [Annual Report on Form 10-K for the year ended December 31, 2000](#).

* the transportation of natural gas through pipelines to markets throughout the United States;

* the generation, transmission and distribution of electricity to markets in the northwestern United States;

¹The term *corpus* is used in natural language processing to denote a collection of related text.

²See, e.g., [Advisory Committee, ESI Checklist, ESI Guidelines, keyword limitations, Sedona Conference](#), and [The Federal Rules of Civil Procedure](#).

³See Krivitsky P, Handcock M (2018). *latentnet: Latent Position and Cluster Models for Statistical Networks*. The Statnet Project <http://www.statnet.org>. R package version 2.9.0, <https://CRAN.R-project.org/package=latentnet> and Krivitsky PN, Handcock MS (2008). "Fitting position latent cluster models for social networks with latentnet." *Journal of Statistical Software*, 24(5).

- * the marketing of natural gas, electricity and other commodities and related risk management and finance services worldwide;
- * the development, construction and operation of power plants, pipelines and other energy related assets worldwide;
- * the delivery and management of energy commodities and capabilities to end-use retail customers in the industrial and commercial business sectors; and
- * the development of an intelligent network platform to provide bandwidth management services and the delivery of high bandwidth communication applications.

As of December 31, 2000, Enron employed approximately 20,600 persons.

For the year ended December 31, 2000, it had operating revenues of \$100,789 million, according to the same report, in which it described one of its businesses as

Enron purchases, markets and delivers natural gas, electricity and other commodities in North America. Customers include independent oil and gas producers, energy- intensive industries, public and investor-owned utility power companies, small independent power producers and local distribution companies. Enron also offers a broad range of price, risk management and financing services including forward contracts, swap agreements and other contractual commitments. Enron's strategy is to enhance the scale, scope, flexibility and speed of its North American energy businesses through developing and acquiring selective assets, securing contractual access to third party assets, forming alliances with customers and utilizing technology such as EnronOnline. With increased liquidity in the marketplace and the success of EnronOnline, Enron believes that it no longer needs to own the same level of physical assets, instead utilizing contracting and market-making activities.

On December 2, 2001, Enron filed for [bankruptcy protection](#).

In less than a year, Enron underwent a complete reversal of fortune as its business strategies ran afoul of applicable regulations, among which were those of the Federal Energy Regulatory Commission (**FERC**).

FERC [became aware](#) of irregularities in the California wholesale electricity market prices, a business in which Enron participated. An orientation to the issues is provided by [testimony](#) before FERC, which provides a concise summary.⁴

Following Enron's bankruptcy, FERC intensified its investigation, including examining the email records of 149 Enron employees. A preliminary [staff report](#) issued six months later.

Motivating Data

FERC obtained approximately 500,000 emails. Copies of these were acquired by Leslie Kaelbling of MIT and [published](#) by William W. Cohen of Carnegie Mellon University. It is one of the largest publicly available datasets of corporate email and is referred to as the Enron Corpus.

⁴The short version, which I can relate as a former California electric utility regulatory official from personal knowledge, is that public electric utilities were losing a large share of industrial customers to self-generation. Many businesses found it cheaper to generate on-site than to pay tariff rates. Foreseeably, residential and business customers without the option to self-generate would come to bear the entire cost of unamortized utility fixed assets (termed *stranded costs*), and rates for retail, commercial and small industrial customers would increase. The adopted solution was to require the utilities to sell their generation plants and buy power on a new public market on a *day-ahead*, tomorrow's estimated demand, and an *hour-ahead* basis for unanticipated demand. Although much thought was devoted to the dangers that participants would game the system to sell at premiums or buy at discounts from market, insufficient consideration was given to multi-participant cooperation.

At the time, electronic record examination (*ediscovery*) in litigation was in a primitive state. It was not uncommon, for example, for paper copies of email to be offered. These would typically be read by teams of freelance attorneys looking for keywords. Advanced technology included scanning with optical character recognition and some proprietary software options to organize emails and capture the status of review.

Much of the focus was directed to keyword searches, sometimes called the *smoking gun* approach. Brute force examination misses opportunities to understand the social networks that reflect how the organization operates, what their concerns are and the haphazard exposure of document reviewers inevitably poses the [Elephant and the Blind Men Problem](#). To triage the corpus quickly and efficiently, it should first be distilled and analyzed in terms of its social network characteristics – who corresponds privately with whom.

Analysis

Data acquisition

I obtained a copy of the [2009 version](#) of the corpus in 2010. It contains copies of emails of a private nature that involve three users who have since requested 27 emails to be [redacted](#). I have removed those.⁵

The following were extracted from the SQL database I prepared for my 2010 analysis on the graph portion of this paper.

body	mediumtext	YES		NULL		
lastword	mediumtext	YES		NULL		
hash	varchar(250)	YES	UNI	NULL		
sender	varchar(250)	YES		NULL		
tos	text	YES		NULL		
mid	varchar(250)	YES		NULL		
ccs	text	YES		NULL		
date	datetime	YES		NULL		
subj	varchar(500)	YES		NULL		
tosctn	mediumint(9)	YES		NULL		
ccsctn	mediumint(9)	YES		NULL		
source	varchar(250)	YES		NULL		

The principal fields used in this paper are:

- sender
- date
- subject
- recipient
- lastword (content in the email that does not occur in its related thread, if any)

Conversion

Each email was a plaintext file⁶ Each user had a directory tree similar to the one below.⁷

Although tedious, traversing the directory tree, parsing the emails and loading them into an SQL database, was accomplished with a combination of Python and Perl scripting and standard bash tools. I do not

⁵Most of my work on data wrangling and preliminary analysis took place in 2010 in Python, relying heavily on the NLTK and networkx packages. For this paper, I did not consult the literature related to graph analysis using the Enron Corpus as an example.

⁶Most had been generated by Microsoft Outlook, but some older emails were produced in IBM Notes, which created some character encoding issues.

⁷This user had 10 directories with 3048 files (the directory tree illustration has been pruned to omit spurious detail) containing 12,147 lines and 69,226 words.




Figure 1: Typical user data

reproduce that process here as it has little bearing on the main topic of this paper.⁸

Data structure

While the emails were not in native format, the plain text versions contained nine principal segments, as shown in the figure below

Deduplication

Using scripting tools, each text file extraction created a *payload* of the new content in the related email, capturing the text between the beginning metadata and the following metadata for email purposes. A *payload* hash, an *md5* encoded message digest⁹ was used in the initial analysis as a primary key to assure the uniqueness of each record. Approximately half of the corpus consisted of duplicates, such as the original message in the sender's sent file and one or more copies in the recipient's inbox, at a minimum. Multiple recipients and recipients who used email folders as a filing system were another source of duplicate messages. Applying this filter reduced the corpus to approximately 250,000 emails.

Text isolation

For natural language processing (NLP) purposes, treating the *payload* rather than the *message body* as the unit of analysis avoided an *echo chamber* effect of *chains* quoting and re-quoting the original message, multiplying the frequency of the words it contained.

Prioritization

Traditional litigation analysis of emails was conducted on the principle that *something may be overlooked*, which delays the value of email in preliminary analysis. Prioritizing always leaves open the option of reviewing the set-asides later.

After deduplication, the first filter applied was to eliminate all email from external addresses that were not also recipients from internal addresses. Spam, newsletters and the like have low information potential. This filter reduced the remaining half of the corpus by half again, leaving approximately 125,000 emails.

A second filter for internal email was used to eliminate broadcast messages and high frequency administrative messages. Indicia of broadcast messages were large numbers of recipients, high frequency, paucity of return correspondence and keyword in context screening. Administrative messages to single recipients were identified by frequency, lack of return correspondence and high frequency words. Many of these were nagging emails concerning the lack of approval of expense reports, for example. This filter reduced the dataset to approximately 35,000.

The third filter limited the dataset to emails sent before Enron's December 2, 2001 bankruptcy. This filter reduced the email count to approximately 13,500, about 2.7% of the original total. A few emails dated "1979-12-31" were reviewed and deleted. The resulting dataset was named *g_enron* for its initial purpose, network graph analysis.

Social network analysis

The nature of social networks

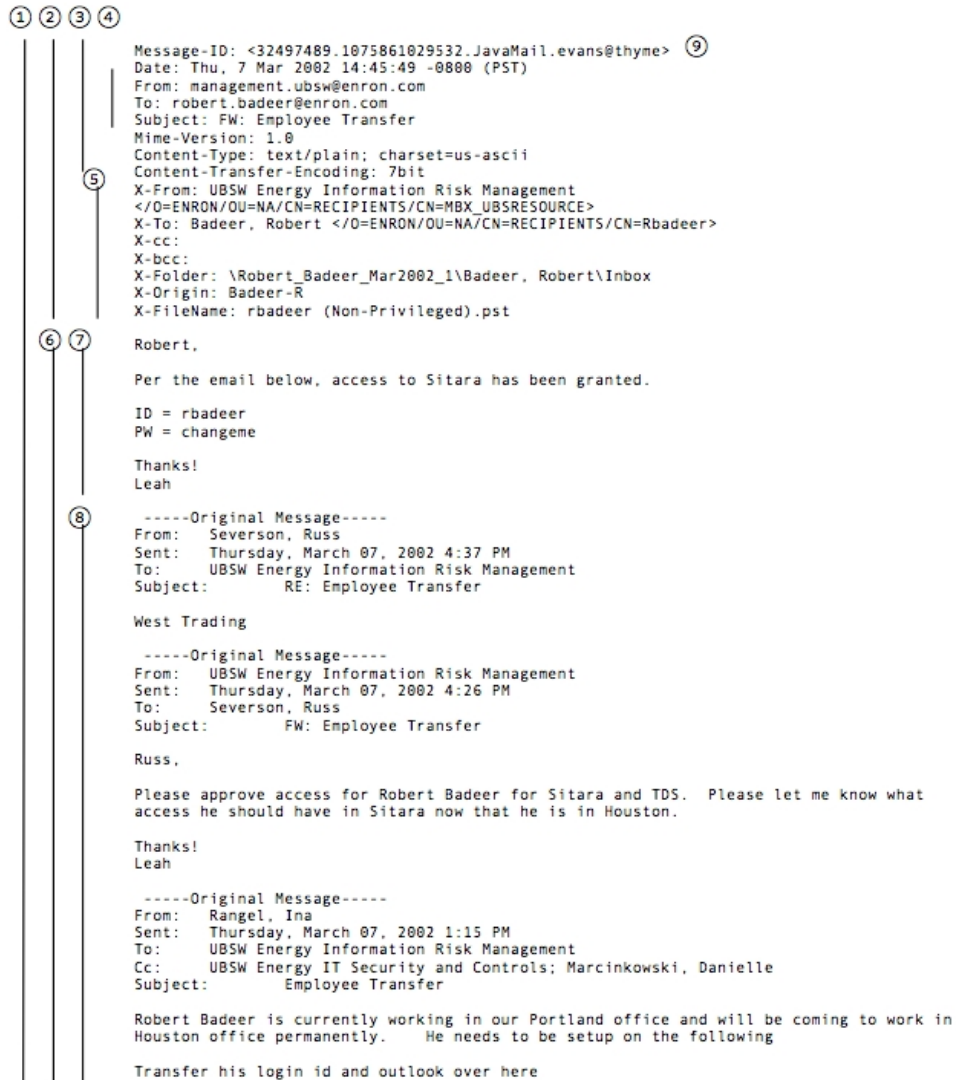
Following the reduction of the corpus, the remaining senders and receivers were natural persons who engaged in mutual correspondence. These constitute *nodes* or *vertices* and their emails *edges*¹⁰. Draw three points and connect them, and you have created three nodes and three edges, a triangle, which is termed a *triad*

⁸For this paper, supplemental processing of the recipient field was necessary and reflected in the script to remove spurious punctuation, such as the newline character embedded as slash-n.

⁹In theory, it is possible that two non-identical sequences of bytes be encoded identically; the probability is low enough to make an *md5* digest usable as a checksum verification, its purpose here.

¹⁰Or *arcs*, when directionality is considered

Parts of an email



- | | |
|---------------------|-------------------------|
| ① Entire email file | ⑦ New content |
| ② Header | ⑧ Chain |
| ③ Visible to user | ⑨ Non-unique identifier |
| ④ Metadata | |
| ⑤ Envelope | |
| ⑥ Message body | |

Figure 2: Structural analysis of an Enron email

graph object. A graph object encapsulates many useful features aside from who knows whom¹¹, including measures of density, centrality, connectedness, separation, clustering and other indicia of how well or poorly embedded in an organization any individual may stand.

Graphs are potentially computationally intensive, which motivated the initial reduction of the selection of emails and users to approximately 1.4% of the emails available for examination. In addition, moving from bigraph directed network to multidirected graph¹² was infeasible.

Graphs are not only a processing unit, they constitute the domain of their own branch of mathematics.¹³

Augmentation and transformation

Each unique Enron address in the reduced dataset was assigned a userid. The primary purpose was to facilitate social network analysis with node identifiers of uniform length; the second, to reduce analyst bias arising from gender stereotyping, frequency of exposure and similar subjective pattern seeking behaviors.

To achieve a computationally practicable dataset for initial social network analysis, emails were limited to single Enron sender to Enron single recipient, reducing the dataset further, to 7,884 emails.

The network composition

Time frame

All emails from January 1, 2000 to December 2, 2001 the date of the [bankruptcy](#) were collected. A handful of messages prior to January 1, 2000 were excluded due to their low counts.

Users

A total of 2,111 unique users are represented. However, all but 1107 users are non-reciprocating or isolated. To identify those, the sender and recipient userids were extracted and converted to a graph object, which will be referred to as the **reduced Enron corpus**. Its attributes are

```
## Network attributes:
## vertices = 56
## directed = TRUE
## hyper = FALSE
## loops = FALSE
## multiple = FALSE
## bipartite = FALSE
## total edges= 532
## missing edges= 0
## non-missing edges= 532
##
## Vertex attribute names:
## sts vertex.names
##
## No edge attributes
```

Definition of terms:

- vertices: users
- directed: from-to and to-from distinguished
- hyper: contains emails from or to multiple users
- loops: includes email from user to herself

¹¹Such as the parlor game [six degrees of Kevin Bacon](#)

¹²A multidirected graph has a single edge to multiple vertices; the analysis is beyond the scope of a term paper for a network as large as the Enron Corpus.


¹³See, e.g., the brief tutorial by [Keijo Ruohonen](#)

- multiple: multi-dimensional object
- bipartite: set of two vertices where no vertex in the same set is connected
- edges: number of emails

The graph can be visualized in several ways. Here, and throughout the paper, a representation based on the Fruchterman-Reingold force-directed algorithm¹⁴ is used to promote visual discrimination.

Graph of reduced Enron corpus


Graph of Enron corpus with isolates



Source: Richard Careaga

Graph of reduced Enron corpus

Graph of Enron corpus without isolates



Source: Richard Careaga

¹⁴Fruchterman, T. M. and Reingold, E. M. (1991), Graph drawing by force-directed placement. Softw. Pract. Exper., 21: 1129-1164. [doi:10.1002/spe.4380211102](https://doi.org/10.1002/spe.4380211102)


Graph objects shown here represent users (vertices) by dots and emails (edges) by lines. The length of the line is not a measure of distance. The visualization algorithm arranges vertices and edges to promote recognition of connections only.

Time series of reduced Enron corpus January 2000-December 2001

Several groups of outliers are apparent, notably mid-May 2001 and the weeks leading up to the [bankruptcy](#).

Time series chart of reduced Enron corpus

January 1, 2000 - December 2, 2001



Source: Richard Careaga

Transitivity

Graph transitivity is a measure of the likelihood that two pairs of vertices (*dyads*) are likely to be strongly connected $A -> B -> C \Rightarrow A -> C$ in the weak form and $A -> B -> C \Leftrightarrow A -> C$ in the strong form. The `sna:gtrans` strong form measure for the graph is 0.9987

User prominence

Graph measures of user prominence

All of the functions described in this section¹⁵, `degree`, `loadcent` and `stresscent` have been run with the `rescale = TRUE` option to normalize them. The functions measure the prominent of a vertex in different ways. The `sna::degree` function relies on measures of incoming and outgoing connections. The `sna::loadcent` function

measures the degree to which a vertex is in a position of brokerage by summing up the fractions of shortest paths between other pairs of vertices that pass through it. [Brandes](#)¹⁶

¹⁵Rejected measures of graph centrality of vertices: betweenness (redundant with `ldctr`); `infocent` (all 1.206801e-13); `closeness` (all 0); `evcent` (asymmetry failure); `bonpow` (system is exactly singular error); `flowbet` (ran without finishing); `graphcent` (all 0)

¹⁶Brandes, U. (2008). "On Variants of Shortest-Path Betweenness Centrality and their Generic Computation." *Social Networks*, 30, 136-145.


The `sna::stresscent` function is a measure of the shortest number of edges that a vertex has to traverse to reach every other vertex in a graph.

Degree

A graph of the top 25 users ranked by `degree` as a sender or receiver is shown below.

Graph of reduced Enron corpus

Graph of Enron corpus after degree filter




Load centrality

A graph of the top 25 users ranked by `load centrality` as a sender or receiver is shown below.

Graph of reduced Enron corpus

Graph of Enron corpus after loadcent filter




Source: Richard Careaga

Stress centrality

A graph of the top 100 users ranked by 'stress centrality users as a sender or receiver is shown below.

Graph of reduced Enron corpus

Graph of Enron corpus, both sender and receiver



Source: Richard Careaga

Usefulness of combination measures

Which of these measures to privilege is not clear. They each present different perspectives of the relative importance of each user in the network, based on different criteria, but none presents an obvious candidate by itself. They are, however, moderately well correlated at high degrees of significance.

Table 1: Pearson’s product-moment correlation: **deg** and **ldctr**

Test statistic	df	P value	Alternative hypothesis	cor
30.71	1105	2.905e-150 * * *	two.sided	0.6786

Table 2: Pearson’s product-moment correlation: **deg** and **sts**

Test statistic	df	P value	Alternative hypothesis	cor
57.47	1105	0 * * *	two.sided	0.8656

Table 3: Pearson’s product-moment correlation: **ldctr** and **sts**

Test statistic	df	P value	Alternative hypothesis	cor
34.26	1105	7.481e-176 * * *	two.sided	0.7177

The union and intersection of the top 25 users using each centrality measure were identified and rejected because their use in subsequent latent network identification either failed or produced unfavorable diagnostics.

Stress centrality was selected because it correlated best with the other two methods and produced satisfactory latent network results as discussed below.

Latent network analysis

Using the top 100 **stresscent** users (rather than the top 25) yields 2,349 edges, each representing an email.

Two latent network analyses were performed. The first selected users who were among the top 100 if they appeared *both* as sender and recipient. The second selected users who were among the top 100 *either* as sender or receiver; that model failed to complete within two hours and was discarded.

The latent cluster random model of senders and receivers

The graph object, prior to modeling, appeared as follows:





Figure 3: Latent graph based on stresscent

Graph of reduced Enron corpus

Graph of Enron corpus, both sender and receiver



Source: Richard Careaga

The `latent::ergmm` model was applied to the graph.

```
c.fit <- ergmm(net_c ~ euclidean(d=2, G=3)+rreceiver,
  control=ergmm.control(store.burnin=TRUE), seed = 2203)
```

The function fits the graph to a latent network model using a Markov chain Monte Carlo algorithm for a Bayesian model fit. The resulting graph visualization identified three clusters. Some vertices show pie slices indicating the relative probabilities of belonging to one of the three clusters. The diagnostics include an intercept estimate, confidence intervals, and a p-value, all of which are satisfactory.

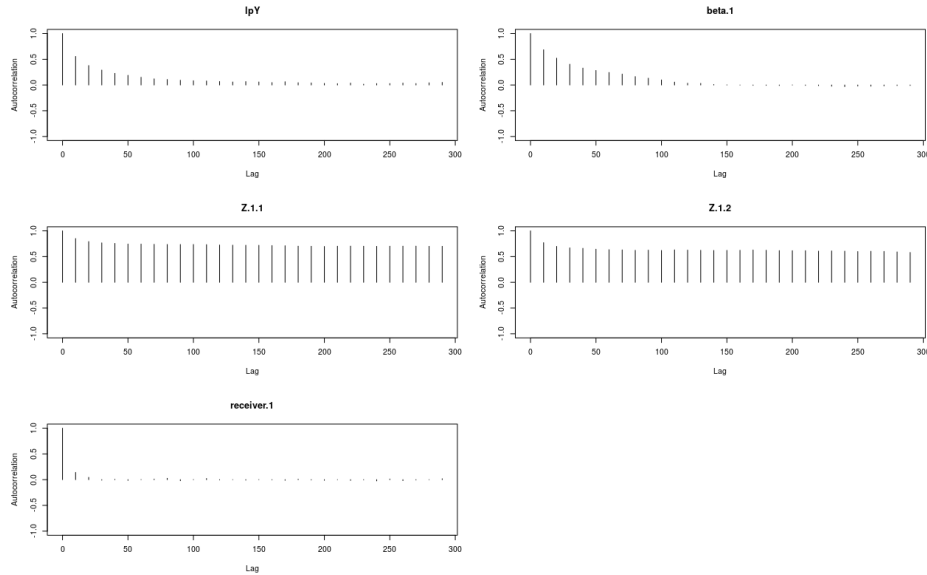


Figure 4: Fit diagnosis part 1

```
##
## =====
## Summary of model fit
## =====
##
## Formula: net_c ~ euclidean(d = 2, G = 3) + rreceiver
## Attribute: edges
## Model: Bernoulli
## MCMC sample of size 4000, draws are 10 iterations apart, after burnin of 10000 iterations.
## Covariate coefficients posterior means:
##      Estimate      2.5%    97.5% 2*min(Pr(>0),Pr(<0))
## (Intercept) -0.71561 -1.07199 -0.3683          < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Receiver effect variance: 0.8183348.
## Overall BIC:          2446.578
## Likelihood BIC:       1617.405
## Latent space/clustering BIC:    638.6051
## Receiver effect BIC:    190.5676
##
## Covariate coefficients MKL:
##      Estimate
## (Intercept) -1.787681
```

Convergence of the model is shown by diagnostic plots of autocorrelation. The log probability (**lpY**) decreases, as does the probability vector (**beta.1**), and the receiver random effect (**receiver1**). The point estimates **Z.1.1** and **Z.1.2** are consistent across lags. The following traces and densities unskewed distributions, and the goodness of fit plots are reasonable.

Goodness of fit diagnostics for in-degree, out-degree and geodesic distance are provided in tabular format and plots. Some excursions in each of the plots appear, indicating the potential benefit of further model tuning.

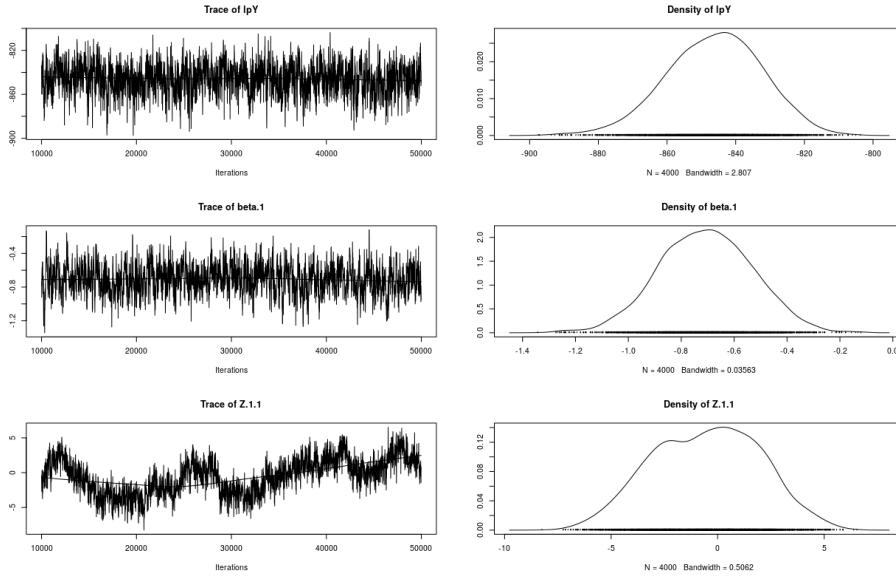


Figure 5: Fit diagnosis part 2

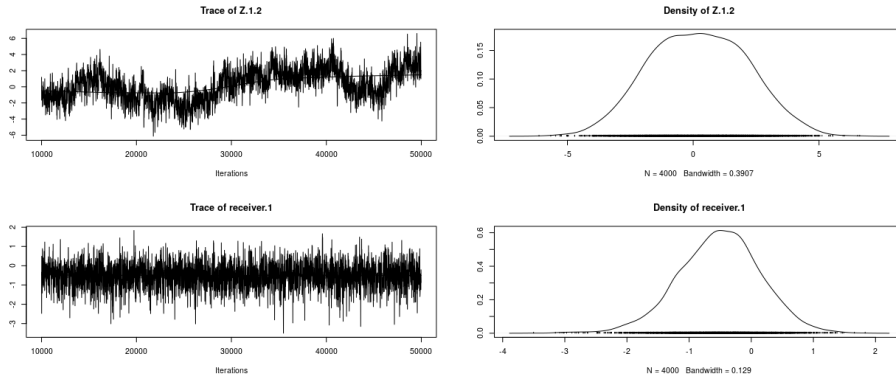


Figure 6: Fit diagnosis part 3

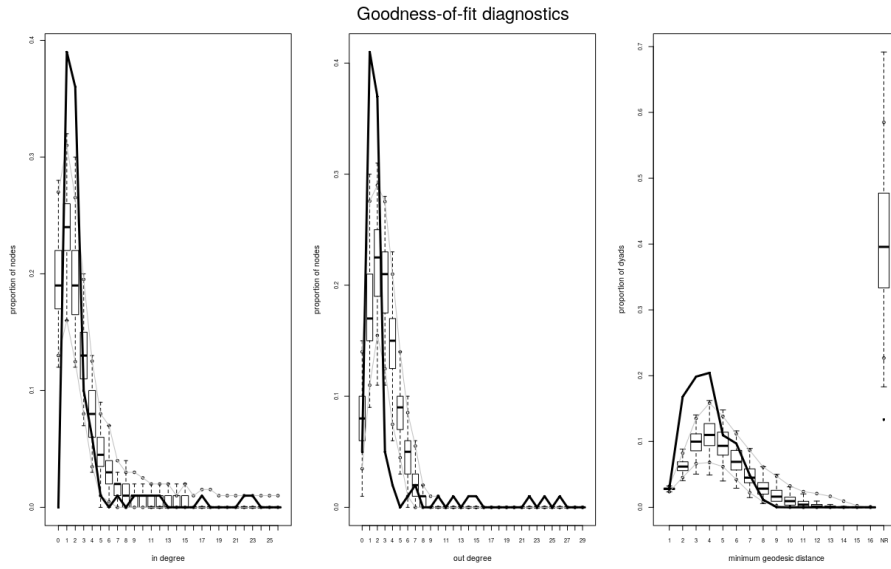


Figure 7: Goodness of fit part 1

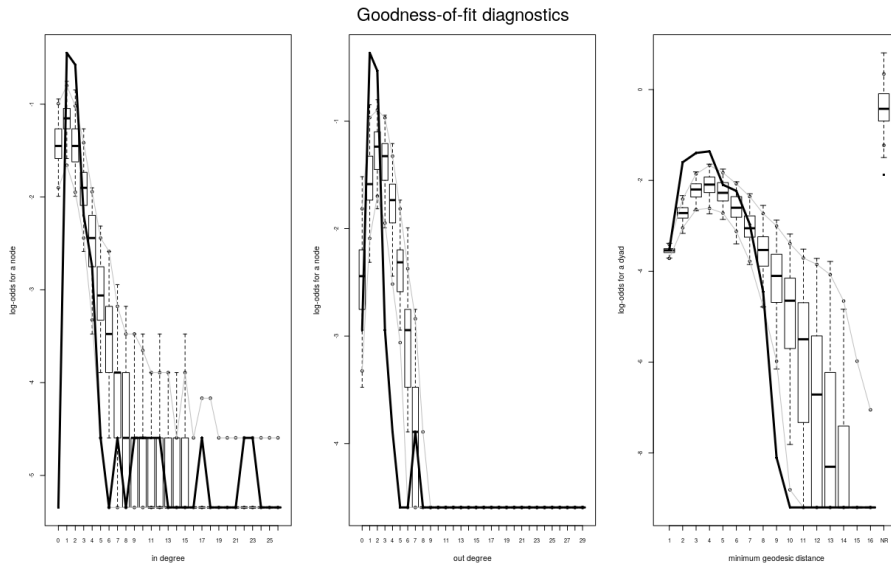


Figure 8: Goodness of fit part 2


```

##
## Goodness-of-fit for in-degree
##
##      obs min  mean max MC p-value
## 0      0  10 19.24  30      0.00
## 1     39  15 23.47  37      0.00
## 2     36  11 19.05  29      0.00
## 3     10   5 13.35  22     0.42
## 4      6   2  8.06  16     0.62
## 5      1   0  5.11  12     0.10
## 6      0   0  2.74   9     0.14
## 7      1   0  1.95   6     0.84
## 8      0   0  1.15   4     0.70
## 9      1   0  0.93   5     1.00
## 10     1   0  0.79   4     1.00
## 11     1   0  0.47   2     0.80
## 12     1   0  0.55   3     0.92
## 13     0   0  0.45   2     1.00
## 14     0   0  0.23   3     1.00
## 15     0   0  0.36   2     1.00
## 16     0   0  0.36   2     1.00
## 17     1   0  0.27   3     0.48
## 18     0   0  0.26   2     1.00
## 19     0   0  0.20   2     1.00
## 20     0   0  0.16   1     1.00
## 21     0   0  0.15   1     1.00
## 22     1   0  0.18   1     0.36
## 23     1   0  0.12   1     0.24
## 24     0   0  0.10   1     1.00
## 25     0   0  0.06   1     1.00
## 26     0   0  0.08   1     1.00
## 27     0   0  0.07   1     1.00
## 28     0   0  0.02   1     1.00
## 29     0   0  0.01   1     1.00
## 30     0   0  0.03   1     1.00
## 31     0   0  0.01   1     1.00
## 32     0   0  0.01   1     1.00
## 33     0   0  0.01   1     1.00
##
## Goodness-of-fit for out-degree
##
##      obs min  mean max MC p-value
## 0      5   2  7.83  15     0.48
## 1     41   9 17.19  30     0.00
## 2     37  14 22.04  33     0.00
## 3      5   6 20.38  31     0.00
## 4      2   4 14.76  26     0.00
## 5      0   2  9.69  21     0.00
## 6      1   0  4.67  10     0.08
## 7      2   0  2.27   8     1.00
## 8      0   0  0.78   4     1.00
## 9      0   0  0.29   2     1.00
## 10     1   0  0.08   1     0.16
## 11     0   0  0.01   1     1.00

```

```

## 12  1  0  0.01  1      0.02
## 14  1  0  0.00  0      0.00
## 15  1  0  0.00  0      0.00
## 22  1  0  0.00  0      0.00
## 24  1  0  0.00  0      0.00
## 26  1  0  0.00  0      0.00
##
## Goodness-of-fit for minimum geodesic distance
##
##      obs   min    mean   max MC p-value
## 1      281   233   283.69   340      1.00
## 2     1663   422   633.80   923      0.00
## 3     1966   619  1003.40  1605      0.00
## 4     2023   620  1143.39  1859      0.00
## 5     1085   456  1010.17  1559      0.90
## 6      960   325   754.28  1360      0.36
## 7      486   173   507.62   925      0.96
## 8      114    36   319.27   627      0.18
## 9         3     2   189.81   427      0.02
## 10       0     0   107.07   340      0.04
## 11       0     0    56.90   301      0.16
## 12       0     0    27.04   215      0.56
## 13       0     0    10.87   121      0.88
## 14       0     0     3.48    51      1.00
## 15       0     0     0.97    19      1.00
## 16       0     0     0.21     7      1.00
## 17       0     0     0.04     1      1.00
## Inf  1319  1622 3847.99  6759      0.00

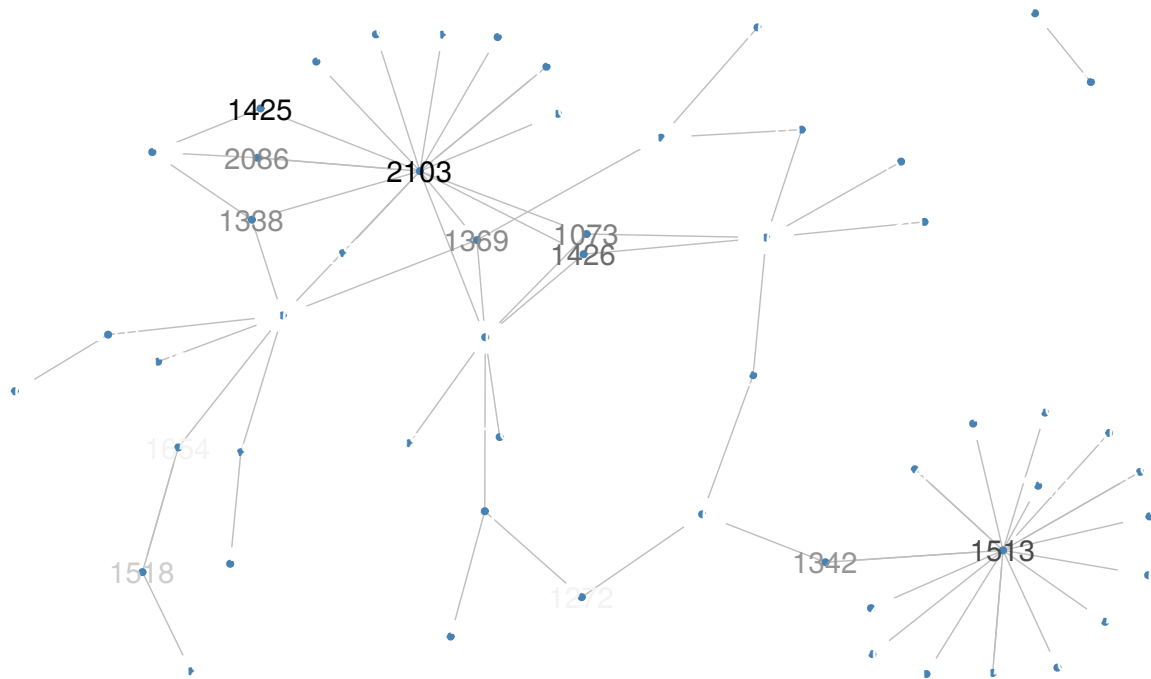
```

The three clusters can be examined separately.¹⁷

¹⁷Highlighted vertices are those included in the top 100 **stresscent** group.

Graph of reduced Enron corpus

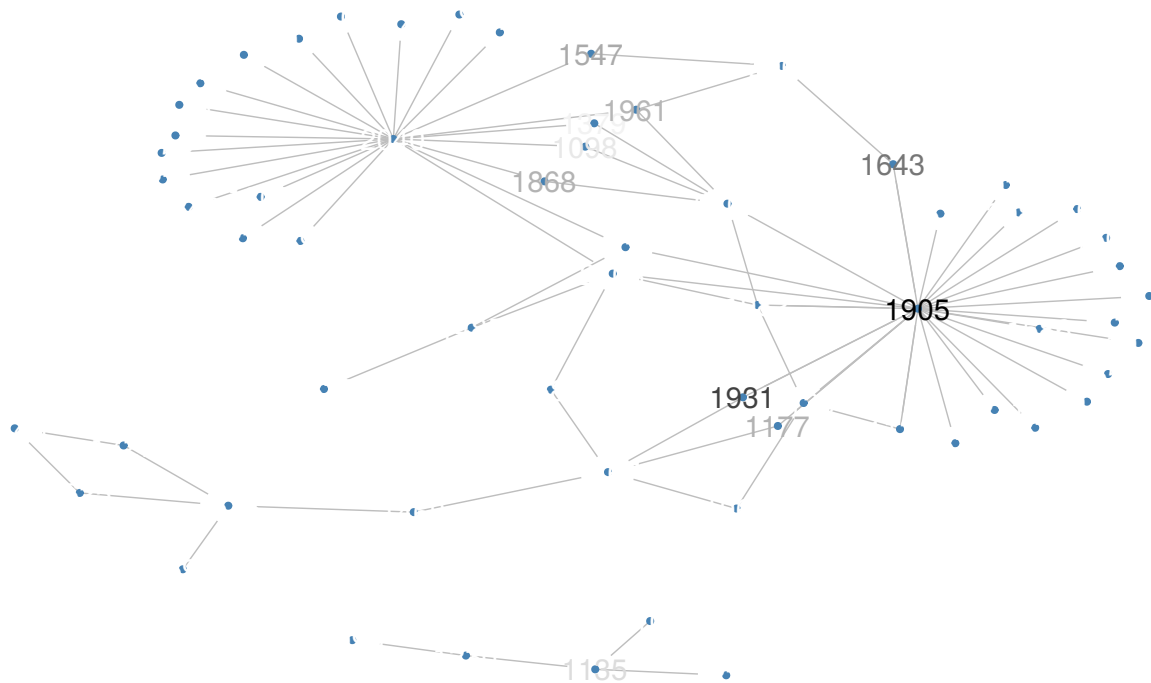
Cluster 1



Source: Richard Careaga

Graph of reduced Enron corpus

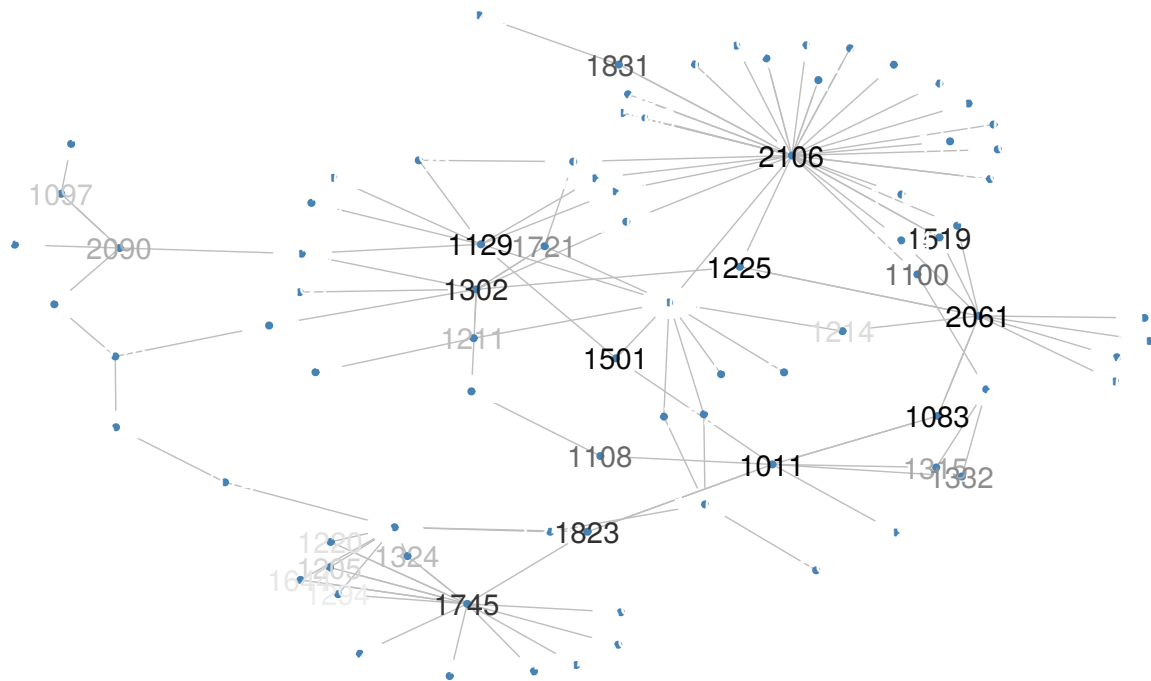
Cluster 2



Source: Richard Careaga

Graph of reduced Enron corpus

Cluster 3



Source: Richard Careaga

Clusters 1 and 3 have the greatest number of vertices and edges. Six distinct userids stand out. Cluster 3 has a relative paucity.

Results

Distinct social networks were identified through latent cluster random graph methods. As a by-product, prominent individuals in the network were identified.

Aside from the visually prominent vertices in the plots above, a simple word frequency analysis of two of the clusters displays markedly differing vocabularies. The third cluster contains no unique terms. strThe clusters are a subset (based on high **stresscent**) score of the larger corpus that has an added field for cluster membership.

Within those clusters are 6,759 distinct words. Of those, 13.36% are unique to Cluster 1; 23.01% are unique to Cluster 2; and 0% are unique to Cluster 3.

Conclusion

The hypothesis of this paper is that social network analysis preprocessing of email text is a feasible method to rapidly identify users who form subgroups with email content of potential interest. Relying solely on metadata (sender/receiver), latent network analysis identified three sub-graphs that have distinct vocabularies.

Credits

Simon Urbanek and Jeffrey Horner (2019). Cairo: R Graphics Device using Cairo Graphics Library for Creating High-Quality Bitmap (PNG, JPEG, TIFF), Vector (PDF, SVG, PostScript) and Display (X11 and Win32) Output. R package version 1.5-10. <https://CRAN.R-project.org/package=Cairo>

Martyn Plummer, Nicky Best, Kate Cowles and Karen Vines (2006). CODA: Convergence Diagnosis and Output Analysis for MCMC, R News, vol 6, 7-11

Winston Chang, (2014). `extrafont`: Tools for using fonts. R package version 0.17. <https://CRAN.R-project.org/package=extrafont>

Francois Briatte (2016). `ggnetwork`: Geometries to Plot Networks with ‘ggplot2’. R package version 0.5.1. <https://CRAN.R-project.org/package=ggnetwork>

Kirill Müller (2017). `here`: A Simpler Way to Find Your Files. R package version 0.1. <https://CRAN.R-project.org/package=here>

Handcock M, Hunter D, Butts C, Goodreau S, Krivitsky P, Morris M (2018). *ergm: Fit, Simulate and Diagnose Exponential-Family Models for Networks*. The Statnet Project <http://www.statnet.org>. R package version 3.9.4, <https://CRAN.R-project.org/package=ergm>

Hunter D, Handcock M, Butts C, Goodreau S, Morris M (2008). “ergm: A Package to Fit, Simulate and Diagnose Exponential-Family Models for Networks.” *Journal of Statistical Software*, 24(3), 1-29.

Barret Schloerke, Jason Crowley, Di Cook, Francois Briatte, Moritz Marbach, Edwin Thoen, Amos Elberg and Joseph Larmarange (2018). `GGally`: Extension to ‘ggplot2’. R package version 1.4.0. <https://CRAN.R-project.org/package=GGally>

Francois Briatte (2016). `ggnetwork`: Geometries to Plot Networks with ‘ggplot2’. R package version 0.5.1. <https://CRAN.R-project.org/package=ggnetwork>

Jeffrey B. Arnold (2019). `ggthemes`: Extra Themes, Scales and Geoms for ‘ggplot2’. R package version 4.1.1. <https://CRAN.R-project.org/package=ggthemes>

Bob Rudis (2019). `hrbrthemes`: Additional Themes, Theme Components and Utilities for ‘ggplot2’. R package version 0.6.0. <https://CRAN.R-project.org/package=hrbrthemes>

Yihui Xie (2019). `knitr`: A General-Purpose Package for Dynamic Report Generation in R. R package version 1.22. Yihui Xie (2015) *Dynamic Documents with R and knitr*. 2nd edition. Chapman and Hall/CRC. ISBN 978-1498716963 Yihui Xie (2014) *knitr: A Comprehensive Tool for Reproducible Research in R*. In Victoria Stodden, Friedrich Leisch and Roger D. Peng, editors, *Implementing Reproducible Computational Research*. Chapman and Hall/CRC. ISBN 978-1466561595

Krivitsky P, Handcock M (2018). *latentnet: Latent Position and Cluster Models for Statistical Networks*. The Statnet Project <http://www.statnet.org>. R package version 2.9.0, <https://CRAN.R-project.org/package=latentnet>

Krivitsky PN, Handcock MS (2008). “Fitting position latent cluster models for social networks with latentnet.” *Journal of Statistical Software*, 24(5).

Butts C (2015). *network: Classes for Relational Data*. The Statnet Project <http://www.statnet.org> R package version 1.13.0.1, <https://CRAN.R-project.org/package=network>

Butts C (2008). “network: a Package for Managing Relational Data in R.” *Journal of Statistical Software*, 24(2). <http://www.jstatsoft.org/v24/i02/paper>

Gergely Daróczi and Roman Tsegelskyi (2018). `pander`: An R ‘Pandoc’ Writer. R package version 0.6.3. <https://CRAN.R-project.org/package=pander>

Carter T. Butts (2016). `sna`: Tools for Social Network Analysis. R package version 2.4. <https://CRAN.R-project.org/package=sna>

Krivitsky P (2019). *statnet.common: Common R Scripts and Utilities Used by the Statnet Project Software*. The Statnet Project <http://www.statnet.org>. R package version 4.2.0, <https://CRAN.R-project.org/package=statnet.common>

Silge J, Robinson D (2016). “tidytext: Text Mining and Analysis Using Tidy Data Principles in R.” *JOSS*, 1(3). doi: 10.21105/joss.00037 <http://doi.org/10.21105/joss.00037>, <http://dx.doi.org/10.21105/joss.00037>

Hadley Wickham (2017). tidyverse: Easily Install and Load the ‘Tidyverse’. R package version 1.2.1.
<https://CRAN.R-project.org/package=tidyverse>

Session Information

```
## R version 3.6.0 (2019-04-26)
## Platform: x86_64-pc-linux-gnu (64-bit)
## Running under: Ubuntu 19.04
##
## Matrix products: default
## BLAS:   /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.8.0
## LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.8.0
##
## locale:
##  [1] LC_CTYPE=en_US.UTF-8      LC_NUMERIC=C
##  [3] LC_TIME=en_US.UTF-8      LC_COLLATE=en_US.UTF-8
##  [5] LC_MONETARY=en_US.UTF-8  LC_MESSAGES=en_US.UTF-8
##  [7] LC_PAPER=en_US.UTF-8     LC_NAME=C
##  [9] LC_ADDRESS=C             LC_TELEPHONE=C
## [11] LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods    base
##
## other attached packages:
##  [1] forcats_0.4.0      stringr_1.4.0      dplyr_0.8.1
##  [4] purrr_0.3.2        readr_1.3.1        tidyr_0.8.3
##  [7] tibble_2.1.2       tidyverse_1.2.1    tidytext_0.2.0
## [10] statnet_2018.10    tsna_0.3.0         ergm.count_3.4.0
## [13] tergm_3.6.0        networkDynamic_0.10.0 sna_2.4
## [16] pander_0.6.3       latentnet_2.9.0    statnet.common_4.3.0
## [19] knitr_1.23         hrbrthemes_0.6.0   ggthemes_4.2.0
## [22] GGally_1.4.0       ergm_3.10.1        network_1.15
## [25] here_0.1           ggnetwork_0.5.1    ggplot2_3.1.1
## [28] extrafont_0.17     Cairo_1.5-10       coda_0.19-2
##
## loaded via a namespace (and not attached):
##  [1] httr_1.4.0          jsonlite_1.6        modelr_0.1.4
##  [4] assertthat_0.2.1    cellranger_1.1.0    yaml_2.2.0
##  [7] robustbase_0.93-5   ggrepel_0.8.1       gdtools_0.1.8
## [10] Rttf2pt1_1.3.7     pillar_1.4.1        backports_1.1.4
## [13] lattice_0.20-38     glue_1.3.1          extrafontdb_1.0
## [16] digest_0.6.19       RColorBrewer_1.1-2  rvest_0.3.4
## [19] colorspace_1.4-1    htmltools_0.3.6     Matrix_1.2-17
## [22] plyr_1.8.4          lpSolve_5.6.13.1    pkgconfig_2.0.2
## [25] broom_0.5.2         haven_2.1.0         mvtnorm_1.0-10
## [28] scales_1.0.0        generics_0.0.2      withr_2.1.2
## [31] lazyeval_0.2.2      cli_1.1.0           readxl_1.3.1
## [34] magrittr_1.5        crayon_1.3.4        evaluate_0.14
## [37] tokenizers_0.2.1    janeaustenr_0.1.5   nlme_3.1-140
## [40] MASS_7.3-51.4       SnowballC_0.6.0     xml2_1.2.0
## [43] tools_3.6.0         hms_0.4.2           trust_0.1-7
## [46] munsell_0.5.0       compiler_3.6.0      rlang_0.3.4
```

## [49]	grid_3.6.0	rstudioapi_0.10	labeling_0.3
## [52]	rmarkdown_1.13	gtable_0.3.0	reshape_0.8.8
## [55]	R6_2.4.0	lubridate_1.7.4	rprojroot_1.3-2
## [58]	stringi_1.4.3	parallel_3.6.0	Rcpp_1.0.1
## [61]	DEoptimR_1.0-8	tidyselect_0.2.5	xfun_0.7

Author contact

Richard Careaga
public@careaga.net
@technocrat
PO Box 3325
Kirkland, WA 98083