

Introduction to social network analysis with R

Pablo Barbera

June 27, 2017

Importing network data into R

In this training session we will be using a small network that indicates interactions in the movie Star Wars Episode IV (<http://evelinag.com/blog/2015/12-15-star-wars-social-network/>). Here, each node is a character and each edge indicates whether they appeared together in a scene of the movie. Edges here are thus *undirected* and they also have weights attached, since they can appear in multiple scenes together.

The first step is to read the list of edges and nodes in this network:

```
edges <- read.csv("data/star-wars-network-edges.csv")
head(edges)
```

```
##      source target weight
## 1      C-3PO  R2-D2      17
## 2       LUKE  R2-D2      13
## 3    OBI-WAN  R2-D2       6
## 4       LEIA  R2-D2       5
## 5        HAN  R2-D2       5
## 6 CHEWBACCA  R2-D2       3
```

```
nodes <- read.csv("data/star-wars-network-nodes.csv")
head(nodes)
```

```
##      name id
## 1    R2-D2  0
## 2 CHEWBACCA  1
## 3    C-3PO  2
## 4     LUKE  3
## 5 DARTH VADER  4
## 6     CAMIE  5
```

For example, we learn that C-3PO and R2-D2 appeared in 17 scenes together.

How do we convert these two datasets into a network object in R? There are multiple packages to work with networks, but the most popular is `igraph` because it's very flexible and easy to do, and in my experience it's much faster and scales well to very large networks. Other packages that you may want to explore are `sna` and `networks`.

Now, how do we create the `igraph` object? We can use the `graph_from_data_frame` function, which takes two arguments: `d`, the data frame with the edge list in the first two columns; and `vertices`, a data frame with node data with the node label in the first column. (Note that `igraph` calls the nodes `vertices`, but it's exactly the same thing.)

```
library(igraph)
g <- graph_from_data_frame(d=edges, vertices=nodes, directed=FALSE)
g
```

```
## IGRAPH UNW- 22 60 --
## + attr: name (v/c), id (v/n), weight (e/n)
## + edges (vertex names):
## [1] R2-D2      --C-3PO      R2-D2      --LUKE
## [3] R2-D2      --OBI-WAN     R2-D2      --LEIA
## [5] R2-D2      --HAN         R2-D2      --CHEWBACCA
## [7] R2-D2      --DODONNA     CHEWBACCA  --OBI-WAN
## [9] CHEWBACCA  --C-3PO      CHEWBACCA  --LUKE
## [11] CHEWBACCA  --HAN        CHEWBACCA  --LEIA
## [13] CHEWBACCA  --DARTH VADER CHEWBACCA  --DODONNA
## [15] LUKE       --CAMIE      CAMIE      --BIGGS
## + ... omitted several edges
```

What does it mean? - U means undirected

- N means named graph
- W means weighted graph
- 22 is the number of nodes
- 60 is the number of edges
- name (v/c) means *name* is a node attribute and it's a character
- weight (e/n) means *weight* is an edge attribute and it's numeric

This is how you access specific elements within the igraph object:

```
V(g) # nodes
```

```
## + 22/22 vertices, named:
## [1] R2-D2      CHEWBACCA  C-3PO      LUKE      DARTH VADER
## [6] CAMIE      BIGGS      LEIA       BERU      OWEN
## [11] OBI-WAN    MOTTI      TARKIN     HAN       GREEDO
## [16] JABBA      DODONNA    GOLD LEADER WEDGE     RED LEADER
## [21] RED TEN    GOLD FIVE
```

```
V(g)$name # names of each node
```

```
## [1] "R2-D2"      "CHEWBACCA" "C-3PO"     "LUKE"     "DARTH VADER"
## [6] "CAMIE"      "BIGGS"      "LEIA"      "BERU"      "OWEN"
## [11] "OBI-WAN"    "MOTTI"      "TARKIN"    "HAN"       "GREEDO"
## [16] "JABBA"      "DODONNA"    "GOLD LEADER" "WEDGE"     "RED LEADER"
## [21] "RED TEN"    "GOLD FIVE"
```

```
vertex_attr(g) # all attributes of the nodes
```

```
## $name
## [1] "R2-D2"      "CHEWBACCA"  "C-3PO"      "LUKE"       "DARTH VADER"
## [6] "CAMIE"      "BIGGS"      "LEIA"       "BERU"       "OWEN"
## [11] "OBI-WAN"    "MOTTI"      "TARKIN"     "HAN"        "GREEDO"
## [16] "JABBA"      "DODONNA"    "GOLD LEADER" "WEDGE"      "RED LEADER"
## [21] "RED TEN"    "GOLD FIVE"
##
## $id
## [1] 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21
```

E(g) # edges

```
## + 60/60 edges (vertex names):
## [1] R2-D2      --C-3PO      R2-D2      --LUKE
## [3] R2-D2      --OBI-WAN    R2-D2      --LEIA
## [5] R2-D2      --HAN        R2-D2      --CHEWBACCA
## [7] R2-D2      --DODONNA    CHEWBACCA  --OBI-WAN
## [9] CHEWBACCA  --C-3PO      CHEWBACCA  --LUKE
## [11] CHEWBACCA  --HAN        CHEWBACCA  --LEIA
## [13] CHEWBACCA  --DARTH VADER CHEWBACCA  --DODONNA
## [15] LUKE        --CAMIE      CAMIE       --BIGGS
## [17] LUKE        --BIGGS      DARTH VADER--LEIA
## [19] LUKE        --BERU      BERU        --OWEN
## + ... omitted several edges
```

E(g)\$weight # weights for each edge

```
## [1] 17 13 6 5 5 3 1 7 5 16 19 11 1 1 2 2 4 1 3 3 2 3 18
## [24] 2 6 17 1 19 6 1 2 1 7 9 26 1 1 6 1 1 13 1 1 1 1 1
## [47] 1 2 1 1 3 3 1 1 3 1 2 1 1 1
```

edge_attr(g) # all attributes of the edges

```
## $weight
## [1] 17 13 6 5 5 3 1 7 5 16 19 11 1 1 2 2 4 1 3 3 2 3 18
## [24] 2 6 17 1 19 6 1 2 1 7 9 26 1 1 6 1 1 13 1 1 1 1 1
## [47] 1 2 1 1 3 3 1 1 3 1 2 1 1 1
```

g[] # adjacency matrix

22 x 22 sparse Matrix of class "dgCMatrix"

```
## [[ suppressing 22 column names 'R2-D2', 'CHEWBACCA', 'C-3PO' ... ]]
```

```
##
## R2-D2      .  3 17 13 . . .  5 . .  6 . .  5 . .  1 . . . . .
## CHEWBACCA  3 .  5 16 1 . . 11 . .  7 . . 19 . .  1 . . . . .
## C-3PO      17 5 . 18 . . 1  6 2 2  6 . .  6 . . . .  1 . .
## LUKE       13 16 18 . . 2 4 17 3 3 19 . . 26 . .  1 1 2 3 1 .
## DARTH VADER .  1 . . . . .  1 . .  1 1 7 . . . . . . . .
## CAMIE      . . .  2 . .  2 . . . . . . . . . . . . . .
## BIGGS      . .  1  4 .  2 .  1 . . . . . . . .  1 2 3 . .

## LEIA       5 11  6 17 1 . 1 .  1 .  1 1 1 13 . . . . . 1 . .
## BERU       . .  2  3 . . .  1 .  3 . . . . . . . . . . . .
## OWEN       . .  2  3 . . .  3 . . . . . . . . . . . . . .
## OBI-WAN    6  7  6 19 1 . .  1 . . . . .  9 . . . . . . .
## MOTTI      . . . .  1 . .  1 . . . .  2 . . . . . . . . .
## TARKIN     . . . .  7 . .  1 . . . .  2 . . . . . . . . .
## HAN        5 19  6 26 . . . 13 . .  9 . . .  1 1 . . . . .
## GREEDO     . . . . . . . . . . . . . .  1 . . . . . . .
## JABBA      . . . . . . . . . . . . . .  1 . . . . . . .
## DODONNA    1  1 .  1 . . . . . . . . . . .  1 1 . . . .
## GOLD LEADER . . .  1 . .  1 . . . . . . . .  1 . 1 1 . .
## WEDGE      . . .  2 . .  2 . . . . . . . .  1 1 . 3 . .
## RED LEADER . .  1  3 . .  3  1 . . . . . . . .  1 3 . 1 .
## RED TEN    . . .  1 . . . . . . . . . . . . . .  1 . .
## GOLD FIVE  . . . . . . . . . . . . . . . . . . . . . .
```

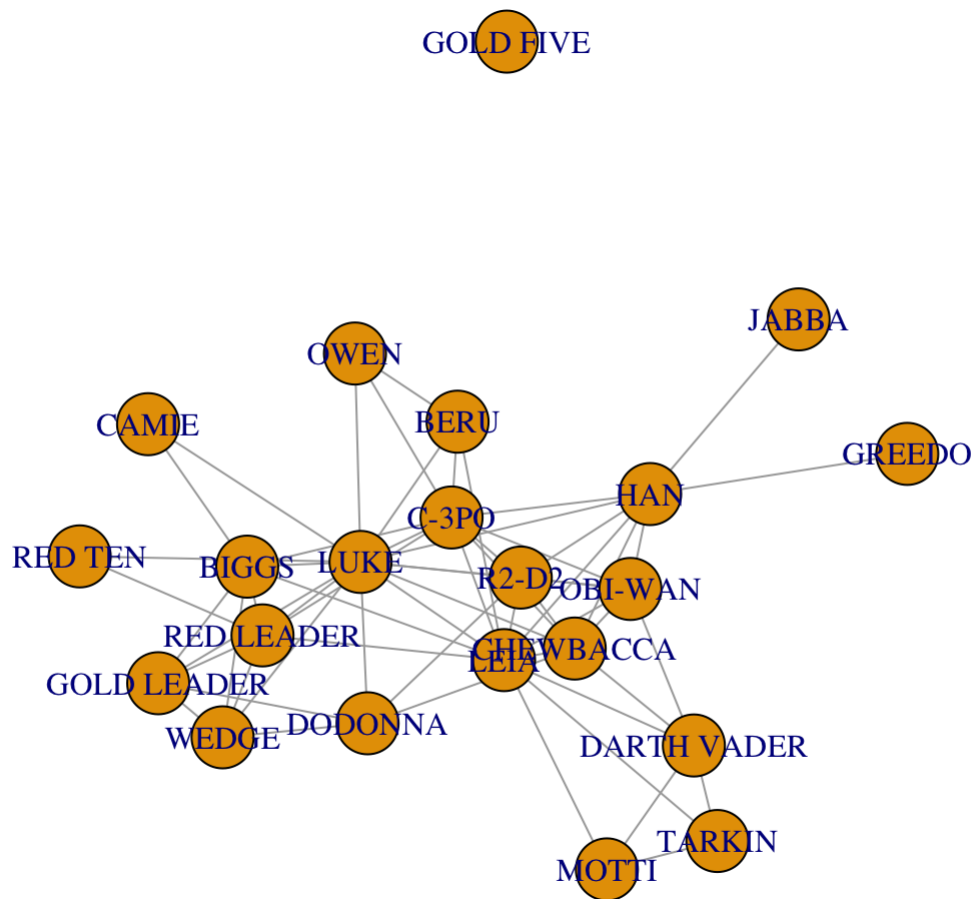
```
g[1,] # first row of adjacency matrix
```

```
##      R2-D2  CHEWBACCA      C-3PO      LUKE DARTH VADER      CAMIE
##      0          3          17          13          0          0
##      BIGGS      LEIA      BERU      OWEN      OBI-WAN      MOTTI
##      0          5          0          0          6          0
##      TARKIN      HAN      GREEDO      JABBA      DODONNA GOLD LEADER
##      0          5          0          0          1          0
##      WEDGE RED LEADER RED TEN  GOLD FIVE
##      0          0          0          0
```

Network visualization

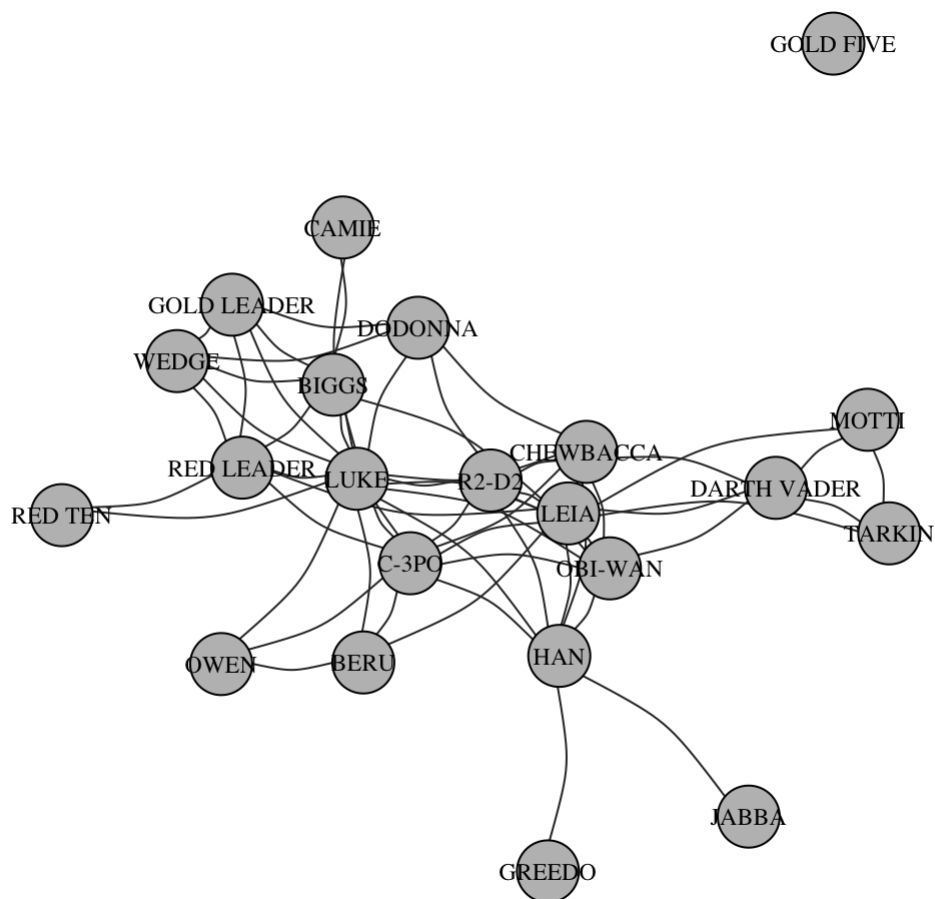
How can we visualize this network? The `plot()` function works out of the box, but the default options are often not ideal:

```
par(mar=c(0,0,0,0))
plot(g)
```




Let's see how we can improve this figure. To see all the available plotting options, you can check `?igraph.plotting`. Let's start by fixing some of these.

```
par(mar=c(0,0,0,0))
plot(g,
      vertex.color = "grey", # change color of nodes
      vertex.label.color = "black", # change color of labels
      vertex.label.cex = .75, # change size of labels to 75% of original size
      edge.curved=.25, # add a 25% curve to the edges
      edge.color="grey20") # change edge color to grey
```

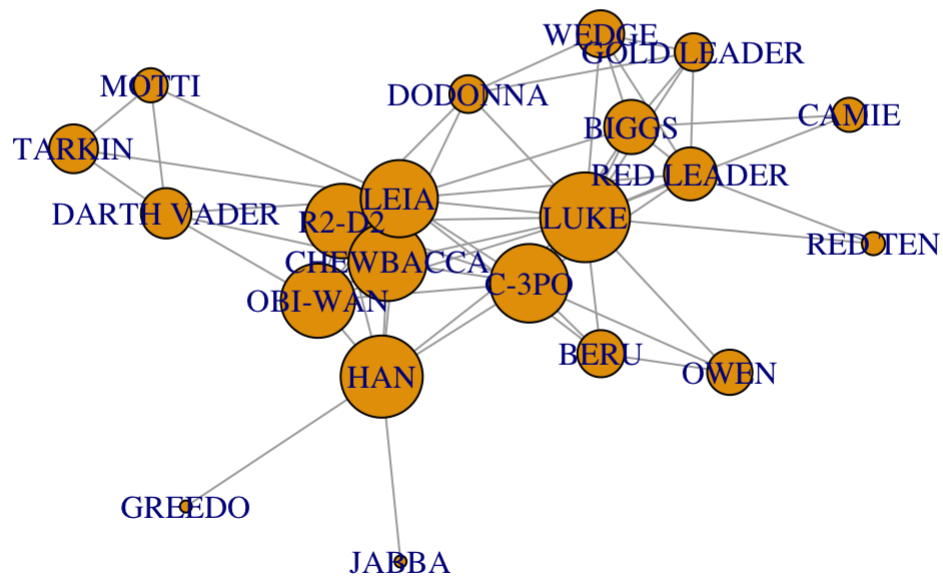


Now imagine that we want to modify some of these plotting attributes so that they are function of network properties. For example, a common adjustment is to change the size of the nodes and node labels so that they match their importance (we'll come back to how to measure that later). Here, `strength` will correspond to the number of scenes they appear in. And we're only going to show the labels of character that appear in 10 or more scenes.


```
V(g)$size <- strength(g)
par(mar=c(0,0,0,0)); plot(g)
```



```
# taking the log to improve it
V(g)$size <- log(strength(g)) * 4 + 3
par(mar=c(0,0,0,0)); plot(g)
```



```
V(g)$label <- ifelse( strength(g)>=10, V(g)$name, NA )
par(mar=c(0,0,0,0)); plot(g)
```

```
# what does `ifelse` do?
nodes$name=="R2-D2"
```

```
## [1] TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [12] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
ifelse(nodes$name=="R2-D2", "yes", "no")
```

```
## [1] "yes" "no" "no" "no" "no" "no" "no" "no" "no" "no" "no" "no"
## [12] "no" "no" "no" "no" "no" "no" "no" "no" "no" "no" "no" "no"
```

```
ifelse(grepl("R", nodes$name), "yes", "no")
```

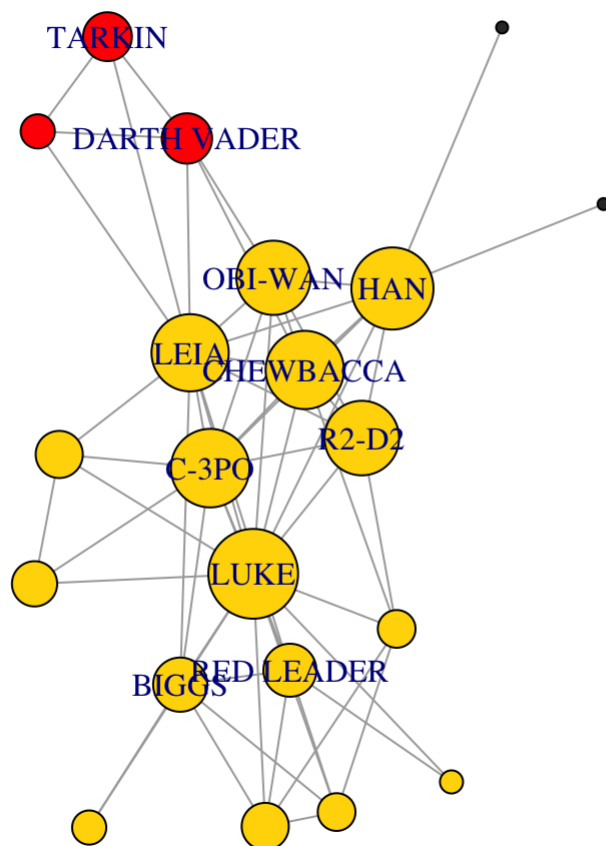
```
## [1] "yes" "no" "no" "no" "yes" "no" "no" "no" "yes" "no" "no" "no"
## [12] "no" "yes" "no" "yes" "no" "no" "yes" "no" "yes" "yes" "no" "no"
```

We can also change the colors of each node based on what side they're in (dark side or light side).

```
# create vectors with characters in each side
dark_side <- c("DARTH VADER", "MOTTI", "TARKIN")
light_side <- c("R2-D2", "CHEWBACCA", "C-3PO", "LUKE", "CAMIE", "BIGGS",
               "LEIA", "BERU", "OWEN", "OBI-WAN", "HAN", "DODONNA",
               "GOLD LEADER", "WEDGE", "RED LEADER", "RED TEN", "GOLD FIVE")
other <- c("GREEDO", "JABBA")
# node we'll create a new color variable as a node property
V(g)$color <- NA
V(g)$color[V(g)$name %in% dark_side] <- "red"
V(g)$color[V(g)$name %in% light_side] <- "gold"
V(g)$color[V(g)$name %in% other] <- "grey20"
vertex_attr(g)
```

```
## $name
## [1] "R2-D2"      "CHEWBACCA"  "C-3PO"      "LUKE"       "DARTH VADER"
## [6] "CAMIE"      "BIGGS"      "LEIA"       "BERU"       "OWEN"
## [11] "OBI-WAN"    "MOTTI"      "TARKIN"     "HAN"        "GREEDO"
## [16] "JABBA"      "DODONNA"    "GOLD LEADER" "WEDGE"      "RED LEADER"
## [21] "RED TEN"    "GOLD FIVE"
##
## $id
## [1] 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21
##
## $size
## [1] 18.648092 19.572539 19.635532 22.439250 12.591581 8.545177 13.556229
## [8] 19.310150 11.788898 11.317766 18.567281 8.545177 12.210340 20.528107
## [15] 3.000000 3.000000 9.437752 9.437752 11.788898 13.259797 5.772589
## [22] -Inf
##
## $label
## [1] "R2-D2"      "CHEWBACCA"  "C-3PO"      "LUKE"       "DARTH VADER"
## [6] NA          "BIGGS"      "LEIA"       NA          NA
## [11] "OBI-WAN"    NA          "TARKIN"     "HAN"        NA
## [16] NA          NA          NA          NA          "RED LEADER"
## [21] NA          NA
##
## $color
## [1] "gold"  "gold"  "gold"  "gold"  "red"   "gold"  "gold"
## [8] "gold"  "gold"  "gold"  "gold"  "red"   "red"   "gold"
## [15] "grey20" "grey20" "gold"  "gold"  "gold"  "gold"  "gold"
## [22] "gold"
```

```
par(mar=c(0,0,0,0)); plot(g)
```



```
# what does %in% do?
1 %in% c(1,2,3,4)
```

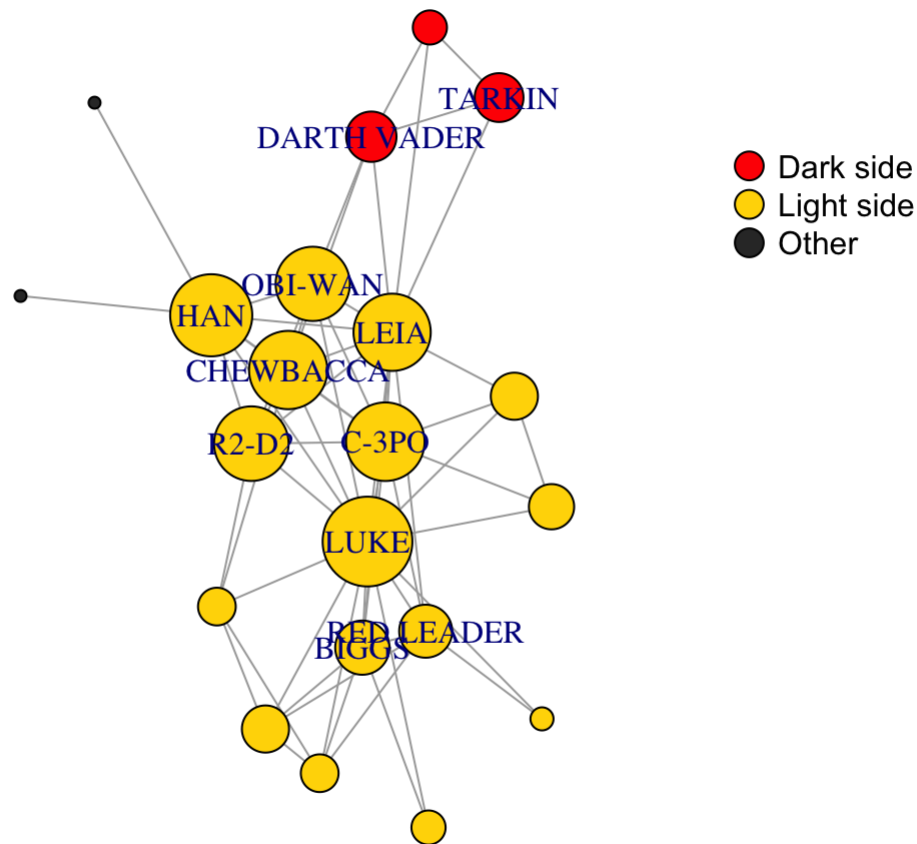
```
## [1] TRUE
```

```
1 %in% c(2,3,4)
```

```
## [1] FALSE
```

If we want to indicate what the colors correspond to, we can add a legend.

```
par(mar=c(0,0,0,0)); plot(g)
legend(x=.75, y=.75, legend=c("Dark side", "Light side", "Other"),
      pch=21, pt.bg=c("red", "gold", "grey20"), pt.cex=2, bty="n")
```




Edge properties can also be modified. For example, here the width of each edge is a function of the log number of scenes those two characters appear together.

```
E(g)$width <- log(E(g)$weight) + 1
edge_attr(g)
```

```
## $weight
## [1] 17 13 6 5 5 3 1 7 5 16 19 11 1 1 2 2 4 1 3 3 2 3 18
## [24] 2 6 17 1 19 6 1 2 1 7 9 26 1 1 6 1 1 13 1 1 1 1 1
## [47] 1 2 1 1 3 3 1 1 3 1 2 1 1 1
##
## $width
## [1] 3.833213 3.564949 2.791759 2.609438 2.609438 2.098612 1.000000
## [8] 2.945910 2.609438 3.772589 3.944439 3.397895 1.000000 1.000000
## [15] 1.693147 1.693147 2.386294 1.000000 2.098612 2.098612 1.693147
## [22] 2.098612 3.890372 1.693147 2.791759 3.833213 1.000000 3.944439
## [29] 2.791759 1.000000 1.693147 1.000000 2.945910 3.197225 4.258097
## [36] 1.000000 1.000000 2.791759 1.000000 1.000000 3.564949 1.000000
## [43] 1.000000 1.000000 1.000000 1.000000 1.000000 1.693147 1.000000
## [50] 1.000000 2.098612 2.098612 1.000000 1.000000 2.098612 1.000000
## [57] 1.693147 1.000000 1.000000 1.000000
```

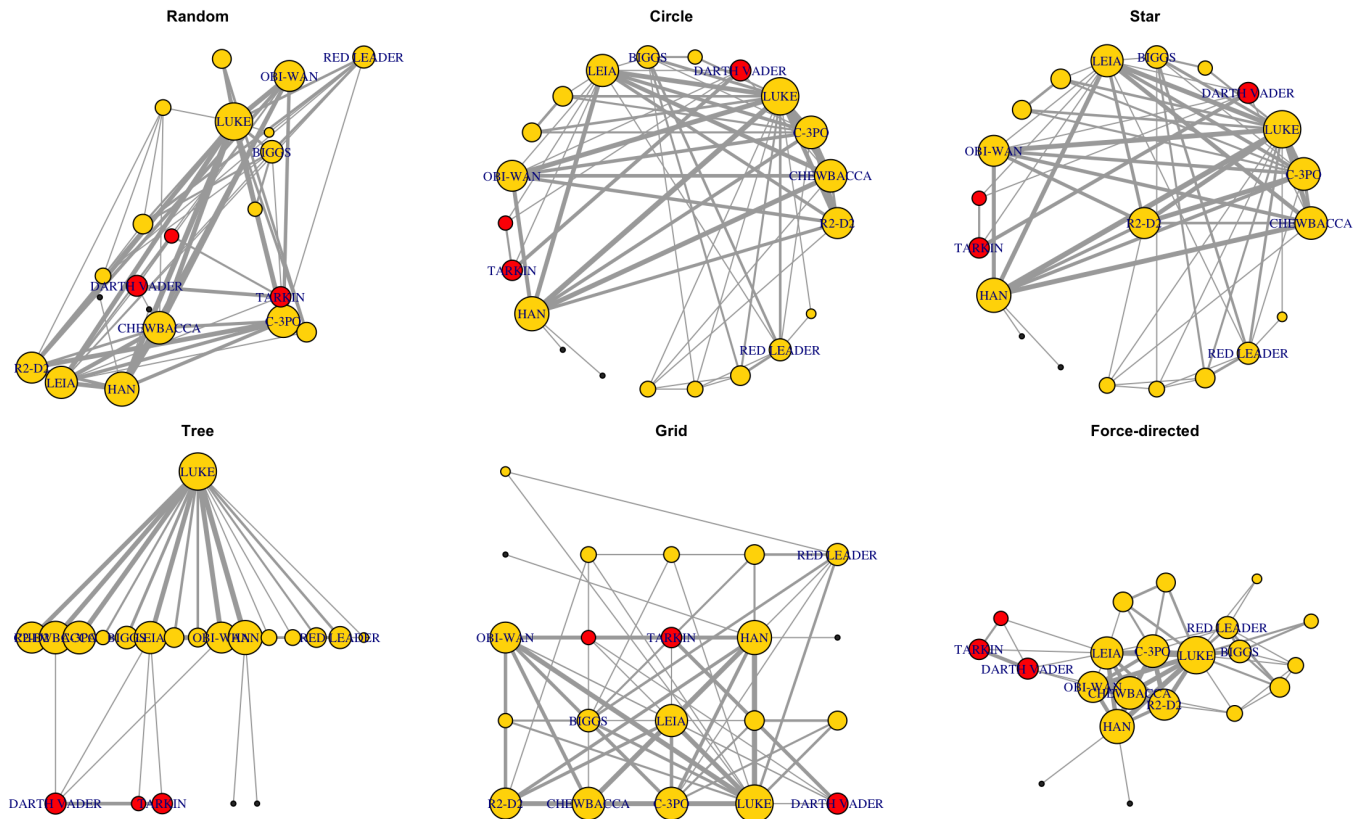
```
par(mar=c(0,0,0,0)); plot(g)
```



Up to now, each time we run the `plot` function, the nodes appear to be in a different location. Why? Because it's running a probabilistic function trying to locate them in the optimal way possible.

However, we can also specify the **layout** for the plot; that is, the (x,y) coordinates where each node will be placed. `igraph` has a few different layouts built-in, that will use different algorithms to find an optimal distribution of nodes. The following code illustrates some of these:

```
par(mfrow=c(2, 3), mar=c(0,0,1,0))
plot(g, layout=layout_randomly, main="Random")
plot(g, layout=layout_in_circle, main="Circle")
plot(g, layout=layout_as_star, main="Star")
plot(g, layout=layout_as_tree, main="Tree")
plot(g, layout=layout_on_grid, main="Grid")
plot(g, layout=layout_with_fr, main="Force-directed")
```



Note that each of these is actually just a matrix of (x,y) locations for each node.

```
l <- layout_randomly(g)
str(l)
```

```
## num [1:22, 1:2] -0.8468 -0.2524 0.1666 -0.0292 0.9343 ...
```

The most popular layouts are force-directed (https://en.wikipedia.org/wiki/Force-directed_graph_drawing). These algorithms, such as Fruchterman-Reingold, try to position the nodes so that the edges have similar length and there are as few crossing edges as possible. The idea is to generate “clean” layouts, where nodes that are closer to each other share more connections in common than those that are located further apart. Note that this is a non-deterministic algorithm: choosing a different seed will generate different layouts.

```
par(mfrow=c(1,2))
set.seed(777)
fr <- layout_with_fr(g, niter=1000)
par(mar=c(0,0,0,0)); plot(g, layout=fr)
set.seed(666)
fr <- layout_with_fr(g, niter=1000)
par(mar=c(0,0,0,0)); plot(g, layout=fr)
```

